

IP Fault Localization Via Risk Modeling

Ramana Rao Kompella[†], Jennifer Yates[‡], Albert Greenberg[‡], and Alex C. Snoeren[†]

[†]*University of California, San Diego*
[†]*{ramana,snoeren}@cs.ucsd.edu*

[‡]*AT&T Labs – Research*
[‡]*{jyates,albert}@research.att.com*

Abstract

Automated, rapid, and effective fault management is a central goal of large operational IP networks. Today’s networks suffer from a wide and volatile set of failure modes, where the underlying fault proves difficult to detect and localize, thereby delaying repair. One of the main challenges stems from operational reality: IP routing and the underlying optical fiber plant are typically described by disparate data models and housed in distinct network management systems.

We introduce a fault-localization methodology based on the use of risk models and an associated troubleshooting system, SCORE (Spatial Correlation Engine), which automatically identifies likely root causes across layers. In particular, we apply SCORE to the problem of localizing link failures in IP and optical networks. In experiments conducted on a tier-1 ISP backbone, SCORE proved remarkably effective at localizing optical link failures using only IP-layer event logs. Moreover, SCORE was often able to automatically uncover inconsistencies in the databases that maintain the critical associations between the IP and optical networks.

1 Introduction

Operational IP networks are intrinsically exposed to a wide variety of faults and impairments. These networks are large, geographically distributed and constantly evolving, with complex hardware and software artifacts. A typical tier-1 network consists of about 1,000 routers from different vendors, with different features, and acting in different roles in the network architecture. Such a network is supported by access and core transport networks, which typically involve at least two orders of magnitude more network elements (optical amplifiers, Dense Wavelength Division Multiplexing (DWDM) systems, ATM/MPLS/Ethernet switches, and so forth). These network elements and associated telemetry generate a large number of management events relating to performance and potential failure conditions.

The essential problem of IP fault management is to monitor the event stream to detect, localize, mitigate and ultimately correct any condition that degrades network

behavior. Unfortunately, operational IP networks today lack intrinsic robustness; serious faults and outages are not infrequent. While existing fault management systems (e.g., [11, 19, 7]) provide great value in automating routine fault management, serious problems can fly “under the radar,” or, once detected, cannot be rapidly localized and diagnosed. To appreciate why this is so, it may help to imagine a network operator faced with the task of IP fault management. After much effort, network hardware has been designed and implemented, the protocols controlling the network have been designed (often in compliance with published standards), and the associated software implemented. In accordance with the network architecture, the network elements have been deployed, connected, and configured. Yet, all these complex endeavors are carried out by multiple teams at rapid pace, involving a large and distributed software component, thus producing operational artifacts far richer in behavior than can ever be approximated in a lab. Errors will be introduced at each stage of network definition and go undetected despite best practices in design, implementation, and testing. External factors, including bugs of all types (memory leaks, inadequate performance separation between processes, etc.) in router software and environmental factors such as DoS attacks and BGP-related traffic events originating in peer networks significantly raise the level of difficulty. It is the task of IP fault management to cope with the result, continually learning and dealing with new failure modes in the field.

In this paper, we introduce a risk modeling methodology that allows for faster, more accurate automatic localization of IP faults to support both real-time and offline analysis. By design, we:

- split our solution into generic algorithmic components (Section 4.2) and problem domain specific components (Section 4.5), and
- create risk models that reflect fundamental architectural elements of the problem domain, but not necessarily implementation details.

As a result, our system is robust to churn in operational networks and is more likely to be extensible to additional system components.

We apply our methodology to the specific problem of fault localization across IP and optical network layers, a difficult problem faced by network operators today. Currently, when IP operations receives router-interface alarms, the systems and staff are often faced with time-intensive manual investigation of what layer the problem occurred in, where, and why. This task is hampered by the architecture of the underlying network: IP uses optics for transport and (in some cases) for self-healing services (e.g., SONET ring restoration) in an overlay fashion. The task of managing each of the two network layers is naturally separated into independent software systems.

Joining dynamic fault data across IP and optical systems is highly challenging—the network elements, supporting standards and information models are totally different. Though there are fields, such as circuit IDs, which can be used to join databases across these systems, automated mechanisms to assure the accuracy of these joins are often limited. Unfortunately, the network elements and protocols provide little help. Path-trace capabilities (counterparts of IP traceroute) are not available in the optical layer, or, if available, do not work in a multi-vendor environment (e.g., where the DWDM systems are provided by multiple vendors). In optical systems such as SONET, there is no counterpart to IP utilization statistics, which might be used to correlate traffic at the IP layer with the optical layer. Both IP and optical network topologies are rapidly changing as equipment is upgraded, network reach is extended, and capacities are re-engineered to manage changing demands.

Our key contribution is the novel and successful application of risk modeling to localize faults across the IP and optical layers in operational networks. Roughly speaking, a physical object such as a fiber span or an optical amplifier represents a *shared risk* for a group of logical entities (such as IP links) at the IP layer. That is, if the optical device fails or degrades, all of the IP components that had relied upon that object fail or degrade. In the literature, these associations are referred to as *Shared Risk Link Groups* or SRLGs [21]. Using only event data gathered at the IP layer, and topology data gathered at both IP and optical layers, we bridge the gap between what the operational information network managers need and what is actually reported at IP layer. Our system, SCORE (Spatial Correlation Engine), relieves operators of the burden of cross-correlating dynamic fault information from two disparate network layers. Once the layer and the location of the fault has been determined, other systems and tools at the appropriate layer can be targeted towards identifying the precise characteristics (for example, rule-based or statistical methods [11, 19]).

2 Troubleshooting using shared risks

Monitoring alarms associated with IP network component failures are typically generated on an individual link basis—for example, a router failure will appear as a failure of all of the links terminating at that router. Best current practice requires a manual correlation of the individual link failure notifications to determine that they are all a result of a common network element (e.g., router). In more complicated failure scenarios, however, it is substantially more challenging to group individual alarms into common groups, and often difficult to even identify in which layer the fault occurred (e.g., in the transport network interconnecting routers, or in the routers themselves). By identifying the set of possible components that could have caused the observed symptoms, shared risk analysis can serve as the first step of diagnosing a network problem. For events being investigated by operations personnel in real time, reducing the time required for troubleshooting directly decreases down time.

2.1 Shared risk in IP networks

Our challenge is to construct a model of risks that represent the set of IP links that would likely be impacted by the failure of each component within the network. The tremendous complexity of the hardware and software upon which an IP network is built implies that constructing a model that accounts for every possible failure mode is impractical. Instead, we identify the key components of the risk model that represent the prevalent network failure modes and those that do not require deep knowledge of each vendor's equipment used within the network. We hasten to add that the better the SRLG modeling of the network, the more precise the fault diagnosis can be. However, as we show later, a solid SRLG model combined with a flexible spatial correlation algorithm can ensure that fault isolation can be robust to missing details in the risk model developed.

The basic IP network topology can be represented as a set of nodes interconnected via links. Inter-domain and intra-domain routing protocols such as OSPF and BGP operate with a basic abstraction of a point-to-point link between routers. Of course, OSPF permits other abstractions such as multi-access and non-broadcast, but a backbone network typically only consists of point-to-point links between routers. Figure 1 illustrates a very simplistic IP network consisting of five nodes connected via six optical links (circuits). Each inter-office IP link is carried on an optical circuit (typically using SONET framing). This optical circuit in turns consists of a series of one or more fibers, optical amplifiers, SONET rings, intelligent optical mesh networks and/or DWDM systems [16]. These systems consist of network elements that provide O-E-O (optical to electrical) conversion and, in the case of SONET rings or mesh optical networks,

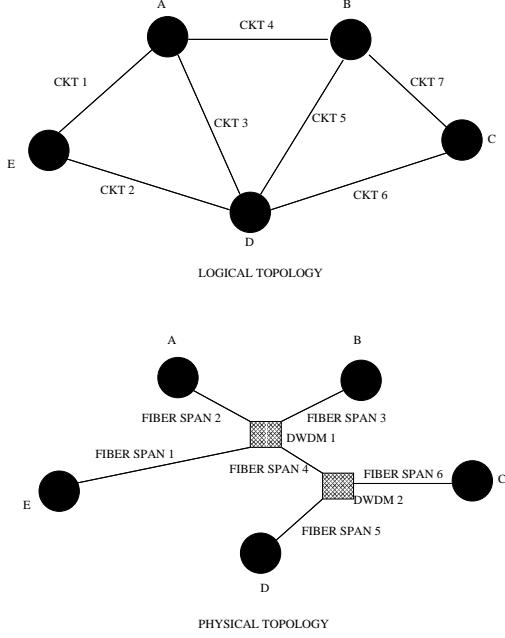


Figure 1: Example illustrating the concept of SRLGs.

protection/restoration to recover from optical layer failures. Multiple optical fibers are then carried in a single conduit, commonly known as a *fiber span*. Typically, each optical component may carry multiple IP links—the failure of these components would result in the failure of all of these IP links. We illustrate this concept in the bottom half of Figure 1, where we show the optical layer topology over which the IP links are routed. In Figure 1, these shared risks are denoted as FIBER SPAN 1 to 6, and DWDM 1 and 2. CKT3 and CKT5 are both routed over FIBER SPAN 4 and thus would both fail with the failure of FIBER SPAN 4. Similarly DWDM 1 is shared between CKT 1, 3, 4 and 5, while CKT 6 and CKT 7 share DWDM 2.

In essence, each network element represents a shared risk among all the links that traverse through this element. Hence, this set of links represents what is known as the Shared Risk Link Group (SRLG), as defined in [21]. This concept is well understood in the context of network planning where backup paths are chosen such that they do not have any SRLG in common with the primary path, and sufficient capacity is planned to survive SRLG failures. However, the application of risk group models to real-time and offline fault analysis has not been well explored.

2.2 Network SRLGs

We now present the shared risk group model that we construct to represent a typical IP network. We divide the model into hardware-related risks and software risks. Note that this model is not exhaustive, and can be ex-

panded to incorporate, for example, additional software protocols.

2.2.1 Hardware-related SRLGs

Fiber: At the lowest level, a single optical fiber carries multiple wavelengths using DWDM. One or more IP links are carried on a given wavelength. All wavelengths that propagate through a fiber form an SRLG with the fiber being the risk element. A single fiber cut can simultaneously induce faults on all of the IP links that ride over that fiber.

Fiber span: In practice, a set of fibers are carried together through a cable. A set of cables are laid out in a conduit. A cut (from, e.g., a backhoe) can simultaneously fail all links carried through the conduit. These set of circuits that ride through the conduit form the fiber span SRLG.

SONET network elements: SONET network elements such as optical amplifiers, add-drop multiplexors etc., are often shared across multiple wavelengths (that represent the circuits). For example, an optical amplifier amplifies all the wavelengths simultaneously – hence a problem in the optical amplifier can potentially disrupt all associated wavelengths. We collectively group these elements together into the SONET network elements group.

Router modules: A router is usually composed of a set of modules, each of which can terminate one or more IP links. A module-related SRLG denotes all of the IP links terminating on the given module, as these would all be subject to failure should the module die.

Router: A router typically terminates a significant number of IP links, all of which would likely be impacted by a router failure (either software or hardware). Hence, all of the IP links terminating on a given router collectively belong to a given router SRLG.

Ports: An individual link can also fail due to the failure of a single port on the router (impacting only the one link), or through other failure modes that impact only the single link. Thus, we also include Port SRLGs in our model. Port SRLGs however are singleton sets consisting of only one circuit. However, we add them in our risk model in order to be able to explain single link failure.

2.2.2 Software-related SRLGs

Autonomous system: An autonomous system (AS) is a logical grouping of routers within the Internet or a single enterprise or provider network (typically managed by a common team and systems). These routers are typically all running a common instance of an intra-domain routing protocol. Although extremely rare, a single intra-domain routing protocol software implementation can cause an entire AS to fail.

OSPF areas: Similar to an AS, an OSPF area is a logical grouping of a set of links for intra-domain routing

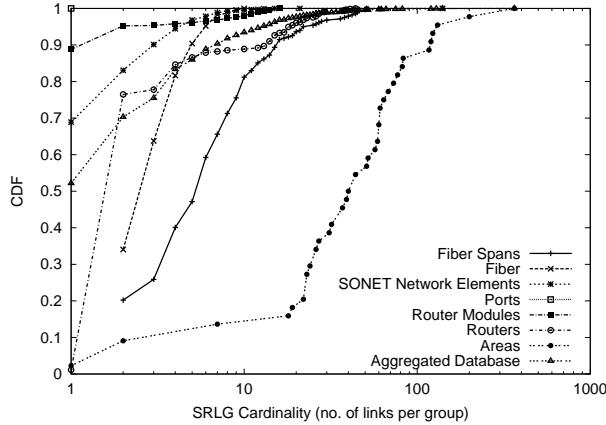


Figure 2: CDF of shared risk among real SRLGs

purposes. It is possible for a faulty routing protocol implementation to cause disruptions across the entire area. Hence, the IP links in a particular area form an OSPF Area SRLG.

Not all SRLGs have corresponding failure diagnosis tools associated with them. For example, a fiber span is a physical piece of conduit that generally cannot indicate to the network operator that it has been cut. Similarly, there is no monitoring at the OSPF area level that can indicate if the whole area was affected. On the other hand, some SONET optical devices can indicate failures in real time. However, these failure indications are usually at wavelength granularity (i.e circuit or link level failures) and hence are not representative of that equipment failure. Diagnosis is therefore based on inference from correlated failures that can be attributed to a particular SRLG. In the absence of fault notifications directly from the equipment, this becomes the only approach to identify the failed component in the network.

2.3 Shared risk in real networks

Spatial correlation is inherently enabled by richness in sharing of risks between links. In particular, spatial correlation will typically be most effective in networks where SRLGs consist of multiple IP links, and each IP link consists of multiple SRLGs. Figure 2 depicts the cumulative distribution of the SRLG cardinality (the number of IP links in each SRLG) in a segment of a large tier-1 IP network backbone (in particular, customer-facing interfaces are not included here). The figure gives an idea of the SRLG cardinality (number of IP links per SRLG) in real networks. We can observe from this figure that, as expected, OSPF areas typically consist of a large number of links (and, hence, are included in their SRLG), whereas port SRLGs (by definition) comprise only a single circuit. In between, we can see that fiber spans typically have a significant number of IP links sharing them,

while SONET network elements typically have fewer. The important observation here is that there is a significant degree of sharing of network components that can be utilized in spatial correlation in real IP networks. Studies of the number of SRLGs along each IP link show similar results. Thus, shared risk group analysis holds great promise for large-scale IP networks.

3 Shared Risk Group analysis

We begin by defining the notation we shall use throughout the remainder of the paper. Define an *observation* as a set of link failures that are potentially correlated, either temporally or otherwise. In other words, if a given set of links fail simultaneously or share a similar pattern or signature of a failure, these events form an observation. A *hypothesis* is a candidate set of circuit failures that could *explain* the observation. That is, a hypothesis is a set of risk groups that contain the set of links seen to fail in a given observation.

The goal, then, of shared risk group analysis is to obtain a hypothesis that best explains a given observation. The principle of Occam’s razor suggests that the simplest explanation is the most likely; hence, we consider the best hypothesis to be the one with the fewest number of risk groups. We note, however, that there could be other formulations of the problem where a best hypothesis is optimizing some other metric.

3.1 Problem formulation

We can define the problem formally as follows. Let $C = \{c_1, c_2, \dots, c_n\}$ denote a set of links, and $G = \{G_1, G_2, \dots, G_m\}$ denote a set of risk groups. Each risk group $G_i \in G$ contains a set of links $G_i = \{c_{i1}, c_{i2}, \dots, c_{ik}\} \subseteq C$ that are likely to fail simultaneously. (We use the terms “links” and “circuits” here to aid intuition, though it will be apparent that the formulation and the algorithm to be described simply deals with sets, and can be applied to arbitrary problem domains). Note that each circuit here can potentially belong to many different risk groups. Given an input observation consisting of events on a subset of circuits, $O = \{c_{e1}, c_{e2}, \dots, c_{em}\}$, the problem is to identify the most probable hypothesis, $H = \{G_{h1}, G_{h2}, \dots, G_{hk}\} \subseteq G$ such that H explains O , i.e., every member of O belongs to at least one member of H and all the members of a given group G_{hi} belong to O . The latter constraint stems from the fact that if a component fails, all the associated member links fail and hence should be a part of the observation. H is a set cover for O ; finding a minimum set cover is known to be NP complete.

The problem can be modeled visually using a bipartite graph as shown in Figure 3. Each circuit, c_i , and group, G_j , is represented by a node in the graph. The bottom partition consists of nodes corresponding to the

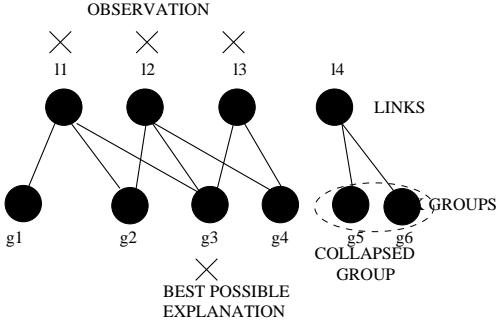


Figure 3: A bipartite graph formulation of the Shared Risk Group problem.

risk groups; the top nodes correspond to circuits. An edge exists between a circuit node and a group node if that circuit is a member of the risk group. Given this bipartite graph and a subset of vertices in the top partition (corresponding to an observation), the problem is to identify the smallest possible set of group nodes that cover the events.

Before proceeding, we observe that if multiple risk groups have the same membership—that is, the same set of circuits may fail for two or more different reasons—it is impossible to distinguish between the causes. We call any such risk groups *aliases*, and collapse all identical groups into one in our set of risk groups. For example, in Figure 3, group g_5 and g_6 have the same membership: l_4 . Hence, g_5 and g_6 are collapsed into a single group as a pre-processing step.

3.2 Greedy approximation

There are potentially many different ways to solve the problem as formulated above; we use a greedy approximation to model imperfect fault notifications and other inconsistencies due to operational realities (as discussed in Section 3.3). Our greedy approximation also reduces the computation cost involved in identifying the most likely hypothesis among all hypotheses (which can potentially be large).

Before presenting the algorithm, however, we must first introduce two metrics we will use to quantify the utility of a risk group. Let $|G_i|$ be the total number of links that belong to the group G_i (known as the cardinality of G_i). Similarly, $|G_i \cap O|$ is the number of elements of G_i that also belong to O . We define the *hit ratio* of the group G_i as $|G_i \cap O|/|G_i|$. In other words, the hit ratio of a group is the fraction of circuits in the group that are part of the observation. The *coverage ratio* of a group G_i is defined as $|G_i \cap O|/|O|$. Basically, the coverage ratio is the portion of the observation explained by a given risk group.

Intuitively, our greedy algorithm attempts to iteratively select the risk group that explains the greatest num-

Algorithm 1 greedyHypothesis(input_links,threshold)

```

1: explained = {} // EmptySet
2: unexplained = input_links;
3: // All groups that contain at least one link
4: groups = getAllGroups(unexplained);
5: while (unexplained ≠ {}) do
6:   // Compute hit and coverage for all groups
7:   hitCoverage(groups, explained, unexplained);
8:   // Find a candidate group for pruning
9:   grp = findCandidateGroup(groups, thresh);
10:  pruneGrp(grp, explained, unexplained);
11:  addGroup(hypothesis, grp);
12: end while
13: return hypothesis;
```

Algorithm 2 findCandidateGroup(groups,threshold)

```

1: for all group such that group.hitratio ≥ threshold do
2:   maxGroup = updateMaxCoverage(group)
3: end for
4: return maxGroup
```

ber of faults in the observation with the least error: in other words, the highest coverage and hit ratios. More concretely, in every iteration, the algorithm computes the hit ratio and coverage ratio for all the groups that contain at least one element of the observation (i.e., the neighborhood of the observation in the bipartite graph). It selects the risk group with maximum coverage (subject to some restrictions on the hit ratio which we shall describe later) and prunes both the group and its member circuits from the graph. In the next iteration, the algorithm recomputes the hit and coverage ratio for the remaining set of groups and circuits. This process repeats, adding the group with the maximum coverage in each iteration to the hypothesis, until finally terminating when there are no circuits remaining in the observation. The pseudocode is presented in Algorithm 1.

The algorithm maintains two separate lists: explained and unexplained circuits. When a risk group is selected for inclusion in the hypothesis, all circuits in the observation that are explained by this risk group are removed from the unexplained list and placed in the explained list. The hit ratio is computed based on the union of the explained and unexplained list, but coverage ratio is computed based only on the unexplained list. The reason for this is straightforward: multiple failures of the same circuit will result in only one failure observation. Hence, the hit ratio of a risk group should not be reduced simply because some other risk group also accounts for the failure observation.

3.3 Modeling imperfections

In our discussion so far, we have skirted the issue of selecting risk groups with hit ratios less than one. What does it mean to have a hypothesis that explains more circuit failures than actually occurred? In a straightforward model, such a result is nonsensical: if the shared component generating the risk group failed, all constituent circuits should have been affected. Operational reality, however, is seemingly contradictory for a number of reasons, including incomplete or erroneous monitoring data, and inaccurate modeling of the shared risk groups.

Although operationally-critical alarms are almost always transmitted using reliable data protocols, some supporting failure-related messages (e.g., syslogs) may be transmitted using unreliable protocols such as UDP. This can result in partial failure observations. In such situations, the accuracy of the diagnosis generated using the lossy messages can be impacted if the data is erroneous or incomplete. For example, consider the failure of a particular optical component, which results in six links failing out of which only five, say, messages make it to the monitoring system. The hit ratio for the risk group representing the shared component is then 5/6. Without expressly allowing for the selection of this risk group, the algorithm would output a hypothesis, that, while plausible, is likely far from reality.

Furthermore, while theoretically it should be possible to precisely model all risk groups, it is impossible in practice to exactly capture all possible failure modes. This difficulty leads to two interesting cases of inaccurate modeling. One is failure to model high-level risk groups (e.g., all links terminating in a particular point of presence may share a power grid) while the other is failure to model low-level risk groups (for example, some internal risk group within a router). Our algorithm needs to be robust against imprecise failure groups and, if possible, learn from real observations. We discuss one real instance of learning of new risk groups from actual failure observations in Section 6.

We allow for these operational realities by allowing risk groups with hit ratios of less than one to be selected as our hypothesis. Specifically, we achieve this by introducing an *error threshold*, and then select the risk group with greatest coverage out of those with hit ratios above this threshold. So, even if a particular circuit is omitted (either due to incorrect modeling or missing data), the error threshold allows consideration of groups that have most links but not quite all and cover a large number of failures.

Note that there could be two different cases once we include groups with hit ratios above a certain error threshold. It is possible that there are in reality only a small number of failures but, as a result of message loss, the algorithm identifies a hypothesis with a larger num-

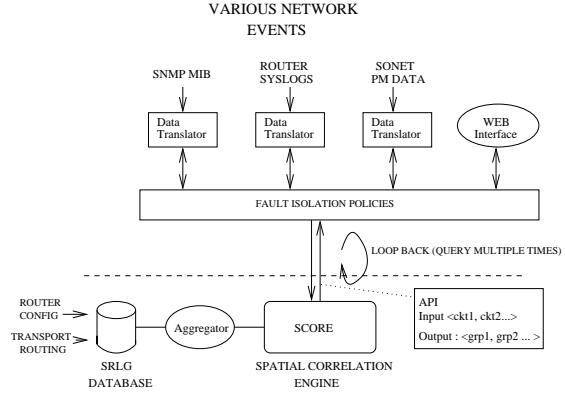


Figure 4: System architecture framework of SCORE.

ber of failures when the error threshold is not used (or error threshold = 1.0). Relaxing the error threshold (≥ 1.0) would account for the message loss, thereby generating a better (smaller) hypothesis. On the other hand, there could be genuinely be a larger number of failures describing our set of observations. In this case, relaxing the error threshold can result in an incorrect hypothesis being generated.

It turns out to be extremely difficult to select a single error threshold for all observations, as it depends greatly on the size of individual risk groups involved in the observation. In practice, we run the algorithm multiple times and generate hypotheses for decreasing error thresholds until a plausible hypothesis is generated. More generally, we can assign a cost function to evaluate the confidence of a particular hypothesis based on the number of component failures in the hypothesis and the threshold used and choose the one with lowest cost.

4 System overview

We created SCORE—the prototype implementation of our system—with generality in mind. Accordingly, key systems and algorithmic components are factored out so that they may be reused in multiple problem domains or in variations for a single problem domain. A stand-alone *spatial correlation module* is driven by an extensible set of problem domain dependent diagnosis processes. Intelligence from the problem domain is built into the SRLG database, and is reflected in the SCORE queries. Figure 4 depicts the SCORE system architecture as it is implemented today. The following subsections describe the various modules in more detail.

4.1 SRLG database

The SRLG database manages relationships between SRLG groups and corresponding links. For example, in our application, the database atoms used to form SRLGs at the SONET layer describe SONET and optical-layer equipment IDs that particular IP links traverse, extracted

from databases populated by operational optical-element management systems. Other risk groups such as areas, routers, modules, etc., are similarly formed from the native databases extracted from the various network elements (e.g., router configurations). We note that the underlying databases track the network and therefore exhibit churn. The SCORE software is currently snapshot driven and copes with churn by reloading multiple times during the course of the day. As mentioned in Section 3 on alias aggregation, we collapse risk groups with identical member links prior to performing spatial correlation.

4.2 Spatial correlation engine

The Spatial Correlation Engine (SCORE) forms the core of the system. This engine periodically loads the spatial database hierarchy and responds to queries for fault localization. SCORE implements the greedy algorithm discussed in Section 3. That is, SCORE obtains the minimum set hypothesis using the SRLG database and a given set of inputs. Optionally, an error threshold can be specified, as described in Section 3.

4.3 Data sources

The set of observations upon which spatial correlation is applied is obtained from network fault notifications and performance reports (including performance-related alarms). These, in turn, come from a wide range of data sources. We discuss below some of the more popular fault- and performance-related data sources that have been used within the SCORE architecture to date. Though we describe certain optical-layer event data sources (such as SONET PM data) and have experimented with such sources in SCORE, only IP event sources were used to obtain the results described in this paper.

4.3.1 IP-layer fault notifications

IP link failures and other faults will be observed by routers and reported to centralized network operations systems via SNMP traps sent from the router. These SNMP traps provide the key event notifications that allow network operators to learn of faults as they occur.

Router operating systems, much like Unix operating systems, log important events as they are observed. These are known as *router syslogs* and provide a wealth of useful information regarding network events. These logs can be used as additional information to complement the SNMP traps and the alarms that they generate. Table 1 shows sample syslog messages for a failure observed on a Cisco router, and another failure observed on an Avici router. The failures are reported at different layers—illustrated here for the SONET, PPP and IP layers (OSPF). Note that there is no standardized format for these messages as they are usually output for debugging purposes.

4.3.2 Performance reports

SNMP performance data is generated by the routers on either a per-interface or per-router basis, as applicable. It typically contains five-minute aggregate measurements of statistics such as traffic volumes, router CPU average utilization, memory utilization of the router, number of packet errors, packet discards, and so on. Performance metrics are also available on a per-circuit basis from SONET network elements along an optical path (as are alarms, although these are not discussed here). Numerous parameters will be reported in, for example, 15-minute aggregates. These include parameters such as coding violations, errored seconds, severely errored seconds (indicative of bit-error rates and outages), and protection switching counts on SONET rings.

4.4 Data translation/normalization

Monitoring data is usually collected from different network elements (routers, SONET DWDM equipment, etc.) and streamed to one or more centralized databases. This various data is usually stored in different formats with independent candidate keys. However, we need to be able to associate the network events (e.g., failures) with the topology information (SRLGs) across layers. To achieve this, we need to use a common key across our SRLG and event notifications.¹

4.5 Fault localization policies

Fault localization is performed on various monitoring data sources (such as those mentioned in the previous section) using flexible data-dependent policies. In Figure 4, fault isolation policies form the bridge between the various monitoring data sources (translators) and the main SCORE engine. These policies dictate how a particular type of fault can be localized. The main functions include:

- **Event clustering.** Clustering events that represent either temporally correlated events or events with similar failure signatures (e.g., spatially correlated)
- **Localization heuristics.** Heuristics that aide in identifying hypotheses that best explain event clusters.

Event clustering: A single failure can result in multiple observations, all close in time, but not necessarily fully synchronized. This may be because the failure was not detected at exactly the same time, the network elements reporting the events may not be all synchronized precisely in time, or propagation delays of the associated messages differ. Thus, data sources that are based on discrete asynchronous events, (e.g., OSPF, syslog messages) need to be clustered together so that individual

¹We map all the databases into link circuit identifiers since the network database itself is organized based on link circuit identifiers. However, any unified format would work equally well.

Syslog Message on Cisco/Avici Routers	Layer	Router
Aug 16 04:01:29.302 EDT: %LINEPROTO-5-UPDOWN: Line protocol on Interface POS0/0, changed state to down	SONET layer	Cisco
Aug 16 04:01:29.305 EDT: %LINK-3-UPDOWN: Interface POS0/0, changed state to down	PPP layer	Cisco
Aug 16 04:01:29.308 EDT: %OSPF-5-ADJCHG: Process 11, Nbr 1.1.1.1 On POS0/0 from FULL to DOWN, Neighbor Down: Interface down or detached	OSPF/IP layer	Cisco
module0036:SUN SEP 12 17:23:29 2004 [030042FF] MINOR:snmp-traps :Sonet link POS 1/0/0 has new adminStatus up and operStatus up.	Sonet Layer	Avici
server0001:SUN SEP 12 17:25:01 2004 [030042FF] MINOR:snmp-traps :PPP link POS 1/0/0 has new adminStatus up and operStatus up.	PPP layer	Avici
server0002:THU AUG 12 07:21:58 2004 [030042FF] MINOR:snmp-traps:OSPF with routerId 1.1.1.1 had non-virtual neighbor state change with neighbor 1.1.1.2 (address less 0) (router id 1.1.1.4) to state Down.	OSPF/IP layer	Avici

Table 1: Syslog messages output by Cisco and Avici routers when a link fails at different layers of the stack. When the link comes back up, the router writes similar messages indicating that each of the layer is back up.

observations (e.g., link down events) become a single observation (multiple-link outage). This clustering captures all the events that are close together in time and assumes that they are potentially related to one or more individual failures.

There are many different ways to cluster events. A naive approach to clustering is based on fixed time bins. For example, we can construct observations (sets of potentially correlated links) by clustering together all events in a fixed five-minute bin. The problem with this approach, however, is the fact that events related to a particular failure can potentially straddle the time-bin boundary. In this case, the quantization will create two different observations for correlated events thus affecting the accuracy of the diagnosis.

In our system, we cluster events based on the length of the gaps between them. We consider the longest chain of events that are spaced apart by less than a set threshold (called the quiet period) as potentially correlated events. The intuition here is that events occurring within a time period less than a given threshold (say 30 seconds) are potentially correlated and can be attributed to the same failure. Note, however, that the quiet period needs to be tuned for the particular problem domain. These clustered events are then fed to the SCORE system to obtain a hypothesis that represents the failed network components.

Localization heuristics: Fault localization often requires heuristics that are either derived intuitively or through domain knowledge to make multiple queries to the system with different parameters in order to obtain higher-confidence hypotheses. The SCORE architecture allows for the straightforward implementation of such troubleshooting policies depending on the problem domain. We implemented one such localization heuristic for handling IP link-down events.

We implemented a simple heuristic that queries SCORE with multiple error thresholds (reducing from 1.0 to 0.5) to obtain a number of different hypotheses. We compare the hypotheses obtained using these relax-

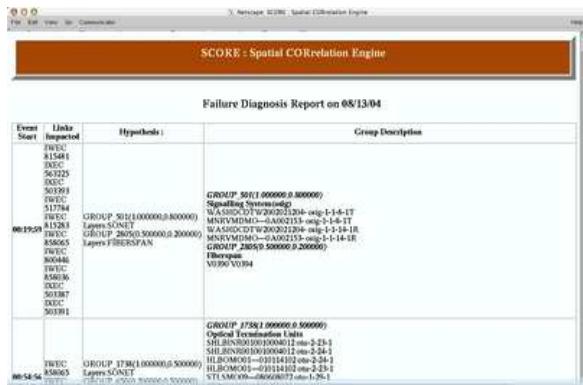


Figure 5: SCORE screen shot.

ations (error thresholds) to account for data inconsistencies or database issues. The most likely hypothesis is selected using a cost function that depends on the error threshold, number of failures in the hypothesis, and, finally, the individual types of groups in the hypothesis. Currently we use the ratio between the number of groups and the threshold; we would like to identify cases where a small relaxation in the threshold (say error threshold of 0.9) can reduce the number of groups significantly.

Another heuristic is to query SCORE using clustered events that have similar signatures. For example, we can further cluster PM data (e.g., packet loss or bit error rates) using additional signatures specific to the data. Links can be grouped together which have similar error counts; these links are more likely to be experiencing the same root cause than links that have very different error counts. This policy is guided by the intuition that correlated events in terms of the actual signature potentially have the same root cause.

4.6 Implementation issues

The core engine that implements the spatial correlation uses the SRLG database, read in periodically, and our al-

gorithm above to identify the most likely cause of each given observation. The main data-structure consists of two hashtables: one for the set of circuits and one for the set of risk groups. Each group consists of the circuit identifiers that can be used to query the circuit's hashtable. This particular implementation allows for fast associations and traversals to implement the spatial correlation algorithm outlined in Section 3. The engine also implements a server to which various diagnosis agents (clients) connect. These clients query the SCORE engine with a set of observations (clusters of link failures). The SCORE engine (server) returns the hypothesis that best explains the observations with which it was queried.

We chose to implement the main SCORE engine in C for efficiency reasons. The total implementation is slightly more than 1000 lines of C code. Obtaining groups from different databases that contain fiber level, fiber-span level, router-level, and other independent views is one of the functions of the SRLG database module. This module is implemented in Perl and consists of approximately 1000 lines of code. The other function of the SRLG database interface is group alias resolution. The group alias resolution algorithm is not a performance bottleneck as it is refreshed fairly infrequently (usually twice a day). This collapsing of risk groups itself is about 200 lines of Perl code. The client module contains the intelligence for the trouble shooting and handles the domain-specific implementations. To achieve flexibility and rapid prototyping, we chose to implement the client module in Perl.

We implemented a Web interface for our SCORE system to conveniently depict the output of our system for operational network data. We used a tabular format for displaying our results, as shown in Figure 5. The Web site displays both current and historical network events, and the corresponding hypotheses output by our system. The first column of the table depicts the observation start and end times as identified using our clustering algorithms. The second column represents the set of links that were impacted during the event (represented here using circuit identifiers). The third and fourth column provide descriptions of the risk groups and associated network components that form the diagnosis report for that observation. The diagnosis report also reports the hit ratio, coverage ratio and the error threshold used for the groups involved in the diagnosis.

5 Simulated faults

We evaluate the performance of SCORE using both artificially generated faults (this section) and real faults (the following section). The goal of our initial experiments is to evaluate the accuracy of the greedy approach within a controlled environment by using emulated faults. We use an SRLG database constructed from the network topol-

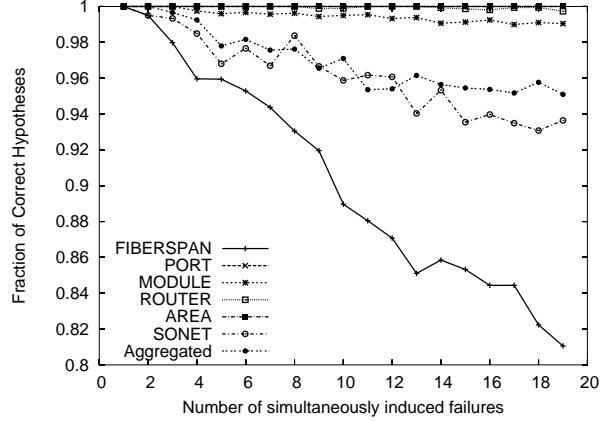


Figure 6: Percentage of correct hypotheses as a function of error probability for various algorithm error thresholds (three simultaneous failures).

ogy and configuration data of a tier-1 service provider's backbone. In our simulation, we inject different numbers of simultaneous faults into the system and evaluate the ability of SCORE to produce the correct hypothesis. We first study the efficacy of the greedy algorithm under ideal operating conditions (no losses in data and no database inconsistencies) before considering noisy data by simulating errors in the SRLG database and event observations.

5.1 Perfect fault notification

To evaluate the accuracy of the SCORE algorithm, we simulated scenarios consisting of multiple simultaneous failures and evaluated the accuracy in terms of the number of correct hypotheses (faults correctly localized by the algorithm) and the number of incorrect hypotheses (those which we did not successfully localize to the correct SRLGs). We randomly generated a given number of simultaneous failures selected from the set of all network risk groups: the set of all SONET components, fiber spans, OSPF areas, routers, and router ports and modules in our SRLG database. Once the faults were selected for a given scenario, we identified the union of all the links that belong to these failures. These link-level failures were input to the SCORE system to generate hypotheses. The resulting hypotheses were then compared with the actual injected failures to determine those which were correctly identified and those which were not.

Figure 6 depicts the fraction of correctly identified hypotheses as a function of the number of simultaneously injected faults, where each data point represents an average across 100 independent simulations. The figure illustrates that the accuracy of the algorithm on these data sets is greater than 99% for ports, modules and routers, irrespective of the number of simultaneous fail-

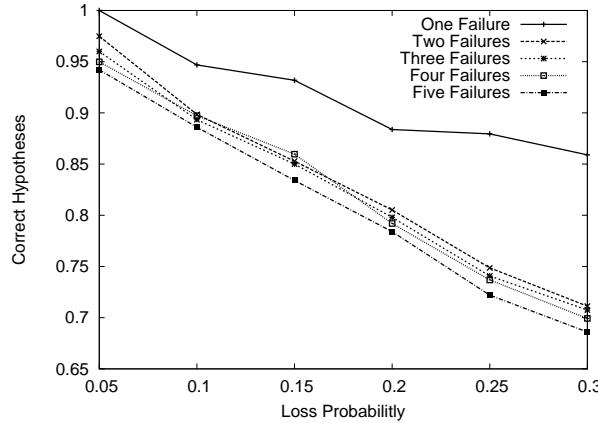


Figure 7: Percentage correct hypotheses as a function of error probability and for varying number of simultaneous failures (error threshold = 0.6).

ures generated. In general, the accuracy of the algorithm decreases as the number of simultaneous failures increases, although the accuracy remains greater than 95% for fewer than five simultaneous failures. In reality, it is unlikely that more than one simultaneous failure will occur (and be reported) at a single point in time. Thus, for failures such as fiber cuts, router failures, and module outages (corresponding to a single simultaneous failure), our results indicate that the accuracy of the system is near 100%. However, it is entirely possible in a large network that multiple independent components will simultaneously be experiencing minor performance degradations, such as error rates, which are reported and investigated on a longer time scale. Thus, the results representing higher number of simultaneous failures are likely indicative of performance troubleshooting. However, we can still conclude that for realistic network SRLGs, the greedy algorithm presented here is highly accurate when we have perfect knowledge of our SRLGs and failure observations.

5.2 Imperfect fault notification

The SRLG model provides a solid, but not perfect representation of the possible failure modes within a complex operational network. Thus, we expect to find scenarios where the set of observations cannot be perfectly described by any SRLG. Similarly, data loss associated with event notifications and database errors are inherent operational realities in managing large-scale IP backbones. In Section 3, we discussed how to adapt the basic greedy algorithm to account for these operational realities. In this section, we evaluate the accuracy of the SCORE algorithm when we have loss in our observations, which may result, for example, from imperfect event notifications (where failures are not reported for

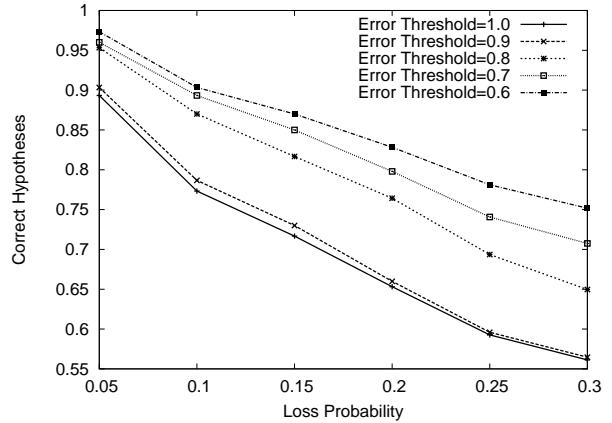


Figure 8: False hypotheses vs error probability (three failures).

whatever reason). We consider three parameters: the error threshold used in the SCORE algorithm, the number of simultaneous failures, and the error probability (which represents the percentage of IP link-failure notifications lost for a given failure scenario).

Figures 7 and 8 demonstrate the accuracy of the algorithm under a range of error probabilities and algorithm error thresholds for a varying number of simultaneous failures. Specifically, the figures plot the percentage of correct hypotheses as a function of the error probability. In Figure 7, the algorithm error threshold is varied from 0.6 to 1.0, while the number of simultaneous failures is set to 3. In Figure 8 the algorithm error threshold is fixed at 0.6 and the number of simultaneous failures is varied from 1 to 5. As expected, increasing the error probability reduces the accuracy of the algorithm. Under three simultaneous failure events and an error probability of 0.1, we can observe from Figure 7 that an algorithm error threshold of between 0.7 and 0.8 restores the accuracy of the SCORE algorithm to around 90%. However, if we mandate perfect matching of failure observations to SRLGs (i.e., error threshold = 1.0), then fault-isolation accuracy drops to around 78%. This shows the necessity and effectiveness of the error thresholds introduced into the algorithm for fault localization in the face of noisy event observation data.

5.3 Performance results

The algorithm's execution time was also evaluated under a range of conditions. In general, the execution time recorded increased as the number of IP links (observations) impacted by the failures increased. This is because all of the SRLGs associated with each of the failed links must be included as part of the candidate set of SRLGs for localization, and thus must be evaluated. Thus, the execution time increased within increasing numbers of

IP link failures, but on average was below 150 ms for up to ten failures. Similarly, the execution time for scenarios involving router failures was typically higher than for other failure scenarios, as the routers typically involved larger numbers of links. Execution times of up to 400 ms were recorded for events involving large routers. However, even in these worst case scenarios, the algorithm is more than fast enough for real-time operational environments.

6 Experience in a tier-1 backbone

The SCORE prototype implementation was recently deployed in a tier-1 backbone network, and used in an offline fashion to isolate IP link failures reported in the network. The implemented system operated on a range of fault and performance data, including IP fault notifications and optical layer performance measures. However, we limit our discussion here to our experience with link failure events reported in router syslogs.

6.1 Localization accuracy

Determining whether or not the SCORE prototype correctly localized a given fault requires identification of the root cause of the fault via other means. In many cases, identifying this root cause involved sifting through large amounts of data and reports—a tedious process at best. We manually confirmed the root cause of 18 faults; we present a comparison with the output reported by the SCORE prototype. We note, however, that our methodology has an inherent bias: we cannot exclude the possibility that there may be a correlation (not necessarily positive) between our ability to diagnose the fault and SCORE’s performance. While it would have been preferable to select a subset of the faults at random, we were not able to easily manually diagnose every fault, nor did we have the resources to consider all faults experienced during SCORE’s deployment.

Table 2 denotes the results of our analysis of each of our 18 faults. For each failure scenario, we report:

- the type of failure that occurred
- a name uniquely identifying the failed component
- the number of SRLG groups localized when the algorithm was run with an error threshold of 1.0
- the threshold used to generate a final conclusion
- the number of SRLGs localized when the algorithm was run with the final threshold
- the number of SRLGs correctly localized
- the number of SRLGs incorrectly localized
- description of the reason why we had to reduce the threshold, or why we were unable to identify a single SRLG as the root cause in certain situations

Overall, we were able to successfully localize all of the faults studied to the SRLGs in which the failed net-

work elements were classified, except where we encountered errors in our SRLG database. However, when we used a threshold of 1.0 (i.e., mandating that an SRLG can be identified if and only if faults were observed on all IP links), then we were typically unsuccessful—particularly for router failures, and for the protocol bug reported. In the majority of the router failures, even though these events corresponded to routers being rebooted, the remote ends of the links terminating at these routers did not always report associated link-level events. This may be due to a number of possible scenarios: the events may never have been logged in the syslogs, data may have been lost from the syslogs, the links may have been operationally shut down and, hence, did not fail at this point in time, or the links were not impacted by the reboot. Independent of why the link notifications were not always observed, the router failures were all successfully localized when the threshold was reduced. This improvement highlights the importance of SCORE’s relaxed approach to localizing faults in operational networks.

Of course, router failures are typically easy to identify through spatial correlation, as all of the links impacted have a common end point (the failed router). However, optical layer impairments can impact seemingly logically independent links at the IP layer if these links are all routed through a common optical component, making them much more difficult to identify.

We studied four different SONET network element failures. The first—an optical amplifier failure—induced faults reported on 13 IP links. Thus, with a threshold of 1.0 our algorithm identified 8 different SRLGs as being involved. However, as the threshold was reduced to 0.9, we were able to isolate the fault to only 2 different SRLGs. Further reductions in this threshold did not, however, further reduce the number of SRLGs to which the fault was localized. Further investigation uncovered an SRLG database problem—where our SONET network element database did not contain any information regarding one of the circuits reporting the fault. Thus, the SCORE algorithm was unable to localize the fault for this particular IP link to the SRLG containing the failed optical amplifier, and instead incorrectly concluded that a router port was also involved (the second SRLG). However, the SRLG containing the failed amplifier was also correctly identified for the other 12 IP links—the lower threshold was required because no fault was observed for one of the IP links routed through the optical amplifier.

This optical amplifier example highlights a particularly important capability of the SCORE system—the ability to highlight potential SRLG database errors. Links missing from databases, incorrect optical layer routing information regarding circuits and other potential errors in databases play havoc with capacity planning and network operations and so must be identified. In this

Type of problem	Component Name	#SRLGS (Thld.=1.0)	Final Thld	#SRLGS (Thld.=Final)	#Correctly localized	#Incorrectly localized	Comment
Router	Router A	27	0.8	1	1	0	No event reported by some links
Router	Router B	20	0.9	3	3	0	No event reported by some links
Router	Router C	12	0.7	1	1	0	No event reported by some links
Router	Router D	1	1	1	1	0	-
Router	Router E	18	0.8	1	1	0	No event reported by some links
Router	Router F	1	1	1	1	0	-
Router	Router G	4	1	4	4	0	-
Module	Module A	1	1	1	1	0	-
Module	Module B	1	1	1	1	0	-
Module	Module C	1	1	1	1	0	-
Optical	Sonet A	8	0.9	2	1	1	No observation reported by one link and database problem
Failed Transceiver	Sonet B	1	1	1	1	0	-
Short term Flap	Sonet C	2	0.7	1	1	0	No observation reported by one link
Optical Amplifier	Sonet D	2	0.6	1	1	0	No observation reported by one link
Fiber Cut	Fiber A	3	0.5	1	1	2	Database problem
Fiber Span	Fiber Span A	1	1	1	1	0	-
Protocol Bug	OSPF Area A	20	0.7	4	4	0	Incorrect SRLG modeling
Protocol Bug	OSPF Area A	4	1	4	4	0	OSPF Area A MPLS enabled interfaces

Table 2: Summary of real, tier-1 backbone failures successfully diagnosed by SCORE.

scenario, the database error was highlighted by the fact that we were unable to identify a single SRLG for a single network failure, even after lowering threshold using in the SCORE algorithm.

The other three SONET failures were all correctly isolated to the SRLG containing the failed network element; in two cases we again had to lower the threshold used within the algorithm to account for links for which we had no failure notification (in one of these cases, the missing link was indeed a result of the interface having been operationally shut down before the failure).

We tested our SCORE prototype on a second, previously identified failure scenario impacted by a SRLG database error (fiber A in table 2). Again, the SCORE system was unable to identify a single SRLG as being the culprit even as the threshold was lowered; no SRLG in the database contained all of the circuits reporting the fault. Again, a database error was highlighted by the inability of the system to map the failure to a single SRLG.

The final case that we evaluated was one in which a low-level protocol implementation problem (commonly known as a software bug) impacted a number of links within a common OSPF area. This scenario occurred over an extended period of time, during which three other independent failures were simultaneously observed in other areas. When a threshold of 1.0 was used in the SCORE algorithm, the event in question was identified as being the result of 20 independent SRLG failures—

a large number even for the extended period of time. As the threshold was reduced to a final value of 0.7, the event was isolated to four individual SRLGs—three SRLGs in other OSPF areas (corresponding to the independent failures) and the OSPF area in question. Thus, the SCORE algorithm was correctly able to identify that the event corresponded to a common OSPF area. However, further investigation uncovered that the reason why not all links in the OSPF area were impacted was that only those interfaces that were currently MPLS-enabled were affected. Thus, an additional SRLG was added to our SRLG database that incorporated the links in a given area that were MPLS enabled—application of this enhanced SRLG database successfully localized all of the SRLGs impacted by the four simultaneous failures with a threshold of 1.0. We conclude that the threshold used in the SCORE algorithm can allow our results to be robust to incomplete modeling of all of the possible SRLGs—any level of modeling of risk groups can be inadequate as there could be more complicated failure scenarios that cannot be modeled by humans perfectly *a priori*. However, SCORE enables operators to continually learn new SRLGs through further analysis of new failure scenarios, thereby enhancing the SRLG models.

6.2 Localization precision

While the 18 faults we studied demonstrate SCORE’s ability to accurately localize faults, it does not give any

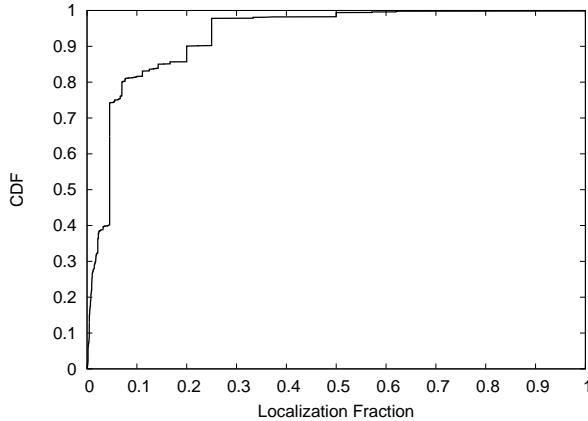


Figure 9: CDF of localization precision out of about 3000 real faults we have been able to localize.

indication of how precise the localization is. Each SRLG to which we localize a fault actually consists of one or more physical components, so isolating a fault to a particular SRLG does not fully isolate the fault to an individual component. In this section, we evaluate the precision of SCORE using a metric we call the *localization precision*. The localization precision of a given observation is defined as the ratio of the number of suspect components after localization to that before localization. In other words, it is the fraction of components that are likely to explain a particular fault (or observation) using our localization algorithm out of all the components that can cause a given fault.

In Figure 9, the cumulative distribution function of the localization precision is shown. From Figure 9, we observe that SCORE localizes faults to less than 5% for more than 40% of the failures and to less than 10% for more than 80% of the failures. We conclude that SCORE identifies likely root causes very precisely from a large set of possible causes for a given failure. We note, however, that we do not know the root cause of all 3000 faults shown here; we cannot speculate on SCORE’s accuracy except for the 18 faults discussed previously.

7 Related work

Network engineers commonly employ the concept of Shared Risk Link Groups to disjoint paths in optical networks; SLRGs are key input into many traffic-engineering mechanisms and protocols such as Generalized Multi-Protocol Label Switching (GMPLS). Due to their importance, recent work has attempted to automatically infer SRLGs [18]. To the best of our knowledge, however, we are the first to use SRLGs in combination with IP-layer fault notifications to isolate failures in the optical hardware of a network backbone without the need for physical-layer monitoring.

Monitoring and management is a challenging problem for any large network. It is not surprising, then, that a number of research prototypes [3, 5, 9, 14, 15, 22] and commercial products have been developed to diagnose problems in IP and telephone networks. Commercial network fault management systems such as SMARTS [19], OpenView [11], IMPACT [12], EXCpert [15], and NetFACT [10] provide powerful, generic frameworks for handling fault indicators, particularly diverse SNMP-based [2] measurements and rule-based correlation capabilities. These systems present unified reporting interfaces to operators and other production network management systems.

In general, previous systems correlate alarms from a particular layer in order to isolate problems at that same layer (e.g., route flapping, circuit failure, etc.). For example, Roughan *et. al.* proposed a correlation-based approach to detect forwarding anomalies such as BGP-related failures [17]. Their approach was to detect events of potential interest by correlating multiple data sources, while our approach uses these events to diagnose the root cause(s), which may or may not be at the BGP layer.

The problem of fault isolation is not limited to networking, of course; similar problems exist in any complex system. Regardless of domain, fault detection systems have taken three basic approaches: rule- or model-based reasoning [1, 6, 11], codebook approaches [19, 23], or machine learning (such as Bayesian and Belief Networks [4, 13, 20]). The difficulty with probabilistic or machine-learning approaches is that they are not prescriptive: it is not clear what sets of scenarios they can handle besides the specific training data. Rule-based and codebook systems (otherwise known as “expert systems”) are often even more specific, only being able to diagnose events that are explicitly programmed. Model-based approaches are more general, but require detailed information about the system under test. Dependency-based systems like ours, on the other hand, allow general inference without requiring undue specificity. Indeed, the specific use of dependency graphs for problem diagnosis has been explored before [8], but not in this particular domain.

Our problem as defined in Section 3.1 falls into the more general class of inference problems which include problems in other domains such as traffic matrix estimation, tomography, etc. Hence, techniques applied in these domains could potentially be used to solve this problem. For example, in [24], the authors reduce the problem of traffic matrix estimation to an ill-posed linear inverse problem and apply a regularization technique to estimate the traffic matrix. Our problem can also be solved using matrix inversion using an incidence matrix to model the risks. While these methods can work well with perfect data, it is unclear how to adapt these

techniques to deal with imperfect loss notifications and SRLG database errors. Our greedy approximation, on the other hand, obtains 95% accuracy for a large class of failures (see Section 5.1) while remaining robust to errors in the SLRG database. Hence, the accuracy improvement from exact techniques is likely to be limited.

8 Conclusions

Using our risk modeling methodology, we have developed a system that accurately localizes failures in an IP-over-optical tier-1 backbone. Given a set of IP-layer events occurring within a small time window, our heuristics pinpoint the shared risk (group of IP and/or optical devices) that best explains these events. Given the harsh operational reality of maintaining complex associations between objects in the two networking layers in separate databases, we find that it is necessary to go beyond identifying the single best explanation, and, instead, to generate a set of likely explanations in order to be robust to transient database glitches.

We put forward a simple, threshold-based scheme that suggests promising explanations while admitting inconsistencies in the data supporting these explanations up to a given threshold. We find that not only does our approach increase the accuracy and robustness of fault localization, but it also provides a way to identify problems with the topology database; we are aware of no alternative automated means of detecting such errors. Accurate modeling of shared risks is critical to IP network design. For example, a misidentification of a shared risk might produce a design believed to be resilient to any single SRLG failure but which in fact is not.

Acknowledgements

This paper benefitted greatly from discussions with Jennifer Rexford, Aman Shaikh, Sriram Ramabhadran, Laurie Knox, Keith Marzullo, Stefan Savage, and comments from the anonymous reviewers and our shepherd, Vern Paxson.

References

- [1] S. Brugnoni, G. Bruno, R. Manione, E. Montariolo, E. Paschetta, and L. Sisto. An expert system for real time fault diagnosis of the Italian telecommunications network. In *Proc. IFIP/IEEE International Symposium on Integrated Network Management*, pages 617–628, Apr. 1993.
- [2] J. Case, M. Fedor, M. Schoffstall, and J. Davin. A Simple Network Management Protocol (SNMP). RFC 1157, May 1990.
- [3] C. S. Chao, D. L. Yang, and A. C. Liu. An automated fault diagnosis system using hierarchical reasoning and alarm correlation. *Journal of Network and Systems Management*, 9(2):183–202, June 2001.
- [4] M. Chen, A. Zheng, J. Lloyd, M. I. Jordan, and E. Brewer. Failure diagnosis using decision trees. In *Proc. International Conference on Autonomic Computing*, pages 36–43, May 2004.
- [5] R. H. Deng, A. A. Lazar, and W. Wang. A probabilistic approach to fault diagnosis in linear light-wave networks. In *Proc. IFIP/IEEE International Symposium on Integrated Network Management*, pages 697–708, Apr. 1993.
- [6] G. Forman, M. Jain, M. Mansouri-Samani, J. Martinka, and A. C. Snoeren. Automated whole-system diagnosis of distributed services using model-based reasoning. In *Proc. IFIP/IEEE Workshop on Distributed Systems: Operations and Management*, Oct. 1998.
- [7] GenSym. Integrity. <http://www.gensym.com>.
- [8] B. Gruschke. Integrated event management: Event correlation using dependency graphs. In *Proc. IFIP/IEEE Workshop on Distributed Systems: Operations and Management*, Oct. 1998.
- [9] P. Hong and P. Sen. Incorporating non-deterministic reasoning in managing heterogeneous network. In *Proc. IFIP/IEEE International Symposium on Integrated Network Management*, pages 481–492, Apr. 1991.
- [10] K. Houk, S. B. Calo, and A. J. Finkel. Towards a practical alarm correlation system. In *Proc. IFIP/IEEE International Symposium on Integrated Network Management*, pages 226–237, May 1995.
- [11] HP Technologies Inc. Open View. <http://www.openview.hp.com>.
- [12] G. Jakobson and M. D. Weissman. Alarm correlation. *IEEE Network*, 7(6):52–59, Nov. 1993.
- [13] K. Jobmann, K.-D. Tuchs, H. Wietgrefe, P. Frohlich, W. Nejdel, and S. Steinfeld. Using neural networks for alarm correlation in cellular phone networks. In *Proc. International Workshop on Applications of Neural Networks in Telecommunications*, May 1997.
- [14] G. Liu, A. K. Mok, and E. J. Yang. Composite events for network event correlation. In *Proc. IFIP/IEEE International Symposium on Integrated Network Management*, pages 247–260, May 1999.
- [15] Y. A. Nygate. Event correlation using rule and object based techniques. In *Proc. IFIP/IEEE International Symposium on Integrated Network Management*, pages 278–289, May 1995.
- [16] R. Ramaswami and K. Sivarajan. *Optical Networks: A Practical Perspective*. Academic Press/Morgan Kaufmann, Feb. 1998.
- [17] M. Roughan, T. Griffin, Z. M. Mao, A. Greenberg, and B. Freeman. IP forwarding anomalies and improving their detection using multiple data sources.

- In *Proc. ACM SIGCOMM NeTs Workshop*, pages 307–312, Aug. 2004.
- [18] P. Sebos, J. Yates, D. Rubenstein, and A. Greenberg. Effectiveness of shared risk link group auto-discovery in optical networks. In *Proc. Optical Fiber Communications Conference*, Mar. 2002.
 - [19] SMARTS. InCharge. <http://www.smarts.com>.
 - [20] M. Steinder and A. Sethi. End-to-end service failure diagnosis using belief networks. In *Proc. IEEE/IFIP Network Operations and Management Symposium*, pages 375–390, Apr. 2002.
 - [21] J. Strand, A. L. Chiu, and R. Tkach. Issues for routing in the optical layer. *IEEE Communications*, 32(2):81–87, Feb. 2001.
 - [22] P. Wu, R. Bhatnagar, L. Epshteyn, M. Bhandaru, and S. Zhongwen. Alarm correlation engine (ACE). In *Proc. IEEE Network Operation and Management Symposium*, volume 3, pages 733–742, Feb. 1998.
 - [23] S. A. Yemini, S. Klinger, E. Mozes, Y. Yemini, and D. Ohsie. High speed and robust event correlation. *IEEE Communications*, 34(5):82–90, May 1996.
 - [24] Y. Zhang, M. Roughan, C. Lund, and D. Donoho. An information-theoretic approach to traffic matrix estimation. In *Proc. ACM SIGCOMM*, pages 301–312, Aug. 2003.